

# COMP433 Assignment 3 Theory,

Omar Elmasaoudi 40255123

## Question 1

$$X = \begin{bmatrix} 1 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 & 1 \end{bmatrix}$$

$$F = \begin{bmatrix} 1 & 4 & 1 \\ 0 & 2 & 1 \\ 2 & 4 & 2 \end{bmatrix}.$$

(a)

No padding, stride = 1.

Output size:

$$H_{\text{out}} = 5 - 3 + 1 = 3, \quad W_{\text{out}} = 5 - 3 + 1 = 3.$$

Example for the top-left element:

$$Y_{0,0} = 1 \cdot 1 + 1 \cdot 4 + 1 \cdot 1 + 0 \cdot 0 + 0 \cdot 2 + 1 \cdot 1 + 0 \cdot 2 + 0 \cdot 4 + 0 \cdot 2 = 7.$$

Doing this for all valid positions gives

$$Y^{(a)} = \begin{bmatrix} 7 & 9 & 8 \\ 7 & 9 & 10 \\ 8 & 5 & 8 \end{bmatrix}.$$

(b)

No padding, stride = 2.

Now  $S = 2$ , still no padding. Output size:

$$H_{\text{out}} = \left\lfloor \frac{5-3}{2} + 1 \right\rfloor = 2, \quad W_{\text{out}} = 2.$$

The resulting feature map is

$$Y^{(b)} = \begin{bmatrix} 7 & 8 \\ 8 & 8 \end{bmatrix}.$$

(c)

1-pixel padding, stride = 1.

We pad  $X$  with one layer of zeros around it ( $P = 1, S = 1$ ). New effective input is  $7 \times 7$ , so output size is  $5 \times 5$ :

$$H_{\text{out}} = 5, \quad W_{\text{out}} = 5.$$

The convolution output is

$$Y^{(c)} = \begin{bmatrix} 3 & 5 & 9 & 10 & 6 \\ 5 & 7 & 9 & 8 & 3 \\ 6 & 7 & 9 & 10 & 7 \\ 9 & 8 & 5 & 8 & 6 \\ 8 & 7 & 3 & 7 & 3 \end{bmatrix}.$$

(d)

MaxPooling, kernel size 2, stride 2.

Max pooling is applied directly to the input image  $X$  (no filter). With kernel size 2 and stride 2:

$$H_{\text{out}} = \left\lfloor \frac{5-2}{2} + 1 \right\rfloor = 2, \quad W_{\text{out}} = 2.$$

We take the maximum in each  $2 \times 2$  window:

$$X = \begin{bmatrix} \boxed{1 & 1} & \boxed{1 & 1} \\ \boxed{0 & 0} & \boxed{1 & 1} \\ \hline \boxed{0 & 0} & \boxed{0 & 0} \\ \boxed{1 & 1} & \boxed{1 & 1} \end{bmatrix} \Rightarrow Y^{(d)} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}.$$

(e)

No padding, stride = 1, dilation = 2.

Dilation  $D = 2$  (on a  $3 \times 3$  filter) gives an effective kernel size

$$K_{\text{eff}} = 1 + (K - 1)D = 1 + 2 \cdot 2 = 5,$$

so with no padding and stride 1 on a  $5 \times 5$  image, the output is just a single value ( $1 \times 1$ ).

We sample every other pixel:

$$\text{positions used in } X : \begin{bmatrix} X_{0,0} & X_{0,2} & X_{0,4} \\ X_{2,0} & X_{2,2} & X_{2,4} \\ X_{4,0} & X_{4,2} & X_{4,4} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 0 & 0 \\ 1 & 0 & 1 \end{bmatrix}.$$

Then

$$Y_{0,0}^{(e)} = 1 \cdot 1 + 1 \cdot 4 + 0 \cdot 1 + 0 \cdot 0 + 0 \cdot 2 + 0 \cdot 1 + 1 \cdot 2 + 0 \cdot 4 + 1 \cdot 2 = 9.$$

So

$$Y^{(e)} = [9].$$

## Question 2

Input image size:

$$64 \times 64 \times 3.$$

Architecture:

- Conv1:  $k = 3 \times 3 \times 32$ ,  $p = 1$ ,  $s = 1$
- Conv2:  $k = 7 \times 7 \times 64$ ,  $p = 0$ ,  $s = 1$
- Pooling: MaxPool  $2 \times 2$ , no overlap ( $s = 2$ )
- Conv3:  $k = 3 \times 3 \times 128$ ,  $p = 1$ ,  $s = 1$
- Conv4:  $k = 7 \times 7 \times 256$ ,  $p = 0$ ,  $s = 1$
- Pooling: MaxPool  $2 \times 2$ , no overlap ( $s = 2$ )
- Flatten
- Dense: size = 10

$$H_{\text{out}} = \left\lfloor \frac{H_{\text{in}} + 2p - k}{s} + 1 \right\rfloor, \quad W_{\text{out}} = \left\lfloor \frac{W_{\text{in}} + 2p - k}{s} + 1 \right\rfloor.$$

**Layer-by-layer dimensions (in  $H \times W \times N$ ):**

Input :  $64 \times 64 \times 3$

Conv1 :  $\underbrace{64 \times 64}_{(64+2 \cdot 1 - 3)/1+1} \times 32$

Conv2 :  $\underbrace{58 \times 58}_{(64+0-7)/1+1} \times 64$

Pooling1 :  $\underbrace{29 \times 29}_{\lfloor \frac{58-2}{2} + 1 \rfloor} \times 64$

Conv3 :  $\underbrace{29 \times 29}_{(29+2 \cdot 1 - 3)/1+1} \times 128$

Conv4 :  $\underbrace{23 \times 23}_{(29+0-7)/1+1} \times 256$

Pooling2 :  $\underbrace{11 \times 11}_{\lfloor \frac{23-2}{2} + 1 \rfloor} \times 256$

Flatten :  $1 \times 1 \times (11 \cdot 11 \cdot 256) = 1 \times 1 \times 30976$

Dense :  $1 \times 1 \times 10$ .

So the output volumes for each layer are:

Conv1:	$64 \times 64 \times 32$
Conv2:	$58 \times 58 \times 64$
Pooling1:	$29 \times 29 \times 64$
Conv3:	$29 \times 29 \times 128$
Conv4:	$23 \times 23 \times 256$
Pooling2:	$11 \times 11 \times 256$
Flatten:	$1 \times 1 \times 30,976$
Dense:	$1 \times 1 \times 10$

### Question 3

- Input images:  $32 \times 32 \times 3$  (RGB).
- Conv1:  $3 \times 3$  filters, 16 filters, stride 1, padding 1.
- Conv2:  $3 \times 3$  filters, 16 filters, stride 1, no padding.
- Average pooling:  $2 \times 2$ , stride 2.
- Fully-connected (FC) layer with ReLU.

- (a) Conv1 keeps the spatial resolution because of padding  $P = 1$ :

$$H_{\text{out}} = \frac{32 + 2 \cdot 1 - 3}{1} + 1 = 32, \quad W_{\text{out}} = 32.$$

Depth is 16. Therefore, the number of neurons is

$$\#\text{neurons in Conv1} = 32 \times 32 \times 16 = 16,384.$$

Each neuron corresponds to one spatial location and one filter.

- (b) Each Conv1 filter has size  $3 \times 3$  and spans all 3 input channels, so:

$$\#\text{weights per filter} = 3 \times 3 \times 3 = 27.$$

There are 16 filters, so

$$\#\text{weights in Conv1} = 27 \times 16 = 432.$$

Each filter has one bias term:

$$\#\text{biases in Conv1} = 16.$$

- (c) The input to Conv2 is  $32 \times 32 \times 16$ . With  $K = 3$ ,  $S = 1$ ,  $P = 0$ :

$$H_{\text{out}} = \frac{32 - 3}{1} + 1 = 30, \quad W_{\text{out}} = 30.$$

Depth is again 16, so

$$\#\text{neurons in Conv2} = 30 \times 30 \times 16 = 14,400.$$

(d) Each Conv2 filter:

$$\#\text{weights per filter} = 3 \times 3 \times 16 = 144.$$

With 16 filters:

$$\#\text{weights in Conv2} = 144 \times 16 = 2,304.$$

Again, one bias per filter:

$$\#\text{biases in Conv2} = 16.$$

(e) After Conv2, the feature maps are  $30 \times 30 \times 16$ . Average pooling with  $2 \times 2$  and stride 2:

$$H_{\text{out}} = \left\lfloor \frac{30 - 2}{2} + 1 \right\rfloor = \left\lfloor \frac{28}{2} + 1 \right\rfloor = 15,$$

$$W_{\text{out}} = \left\lfloor \frac{30 - 2}{2} + 1 \right\rfloor = 15.$$

Depth remains 16. So the encoder output volume is

$$15 \times 15 \times 16.$$

Flattening this gives

$$\#\text{features to FC} = 15 \times 15 \times 16 = 3,600.$$

(f) Using multiple convolutional layers with small filters is better because it reduces the number of parameters compared to a single large filter with the same receptive field, which helps prevent overfitting. Additionally stacking layers introduces multiple non-linearities, allowing the network to learn hierarchical features.

## Question 4

- The input to logistic regression is the flattened encoder output, which has dimensionality 3,600.
- For binary classification, logistic regression applies a linear transformation followed by a sigmoid:

$$z = \mathbf{w}^\top \mathbf{h} + b, \quad \hat{y} = \sigma(z),$$

where  $\mathbf{h} \in R^{3600}$  is the flattened CNN output.

So the logistic regression model takes a 3,600-dimensional input.

### Question 5

$$X = \begin{bmatrix} -1 & 0 & 1 \\ 2 & 3 & 4 \\ 5 & 6 & 7 \\ 8 & 9 & 10 \end{bmatrix}, \quad F = \begin{bmatrix} 2 & 1 \\ 0 & -1 \end{bmatrix}.$$

(a) Output size:

$$H_{\text{out}} = 4 - 2 + 1 = 3, \quad W_{\text{out}} = 3 - 2 + 1 = 2.$$

Using cross-correlation:

$$Y_{i,j} = \sum_{u=0}^1 \sum_{v=0}^1 X_{i+u, j+v} F_{u,v}.$$

Computations:

$$Y = \begin{bmatrix} -5 & -3 \\ 1 & 3 \\ 7 & 9 \end{bmatrix}.$$

(b) Apply ReLU( $z$ ) = max(0,  $z$ ) element-wise:

$$Y^{(\text{ReLU})} = \begin{bmatrix} 0 & 0 \\ 1 & 3 \\ 7 & 9 \end{bmatrix}.$$

(c) After  $2 \times 2$  average pooling, stride 2.

With  $H = 3$ ,  $W = 2$ ,  $K = 2$ ,  $S = 2$ , the pooled output is  $1 \times 1$ : we take the top-left  $2 \times 2$  block

$$\begin{bmatrix} 0 & 0 \\ 1 & 3 \end{bmatrix},$$

whose average is

$$\frac{0 + 0 + 1 + 3}{4} = 1.$$

So the final pooled activation map is

$$Y^{(\text{pool})} = [1].$$

### Question 6

We have a simple 1D RNN with:

$$h_t = w_{\text{in}} x_t + w_{\text{hh}} h_{t-1}, \quad y = w_{\text{hy}} h_L,$$

and loss

$$\ell = \frac{1}{2}(y - y_r)^2.$$

(a) By definition (no activation function):

$$h_t = w_{\text{in}}x_t + w_{\text{hh}}h_{t-1}, \quad t = 1, \dots, L.$$

(b) Forward pass for  $x = [0.3, 0.4, 0.2]$ .

Given

$$x = [x_1, x_2, x_3] = [0.3, 0.4, 0.2], \quad h_0 = 0.5, \quad w_{\text{in}} = w_{\text{hh}} = w_{\text{hy}} = 0.2,$$

we get:

$$h_1 = 0.2 \cdot 0.3 + 0.2 \cdot 0.5 = 0.06 + 0.10 = 0.16,$$

$$h_2 = 0.2 \cdot 0.4 + 0.2 \cdot 0.16 = 0.08 + 0.032 = 0.112,$$

$$h_3 = 0.2 \cdot 0.2 + 0.2 \cdot 0.112 = 0.04 + 0.0224 = 0.0624,$$

$$y = w_{\text{hy}}h_3 = 0.2 \cdot 0.0624 = 0.01248.$$

(c) Gradients

Let  $L = 3$ . First note:

$$\frac{\partial \ell}{\partial y} = y - y_r.$$

Since  $y = w_{\text{hy}}h_3$ ,

$$\frac{\partial \ell}{\partial w_{\text{hy}}} = (y - y_r)h_3,$$

$$\frac{\partial \ell}{\partial h_3} = (y - y_r)w_{\text{hy}}.$$

The recurrence is

$$h_1 = w_{\text{in}}x_1 + w_{\text{hh}}h_0, \quad h_2 = w_{\text{in}}x_2 + w_{\text{hh}}h_1, \quad h_3 = w_{\text{in}}x_3 + w_{\text{hh}}h_2.$$

We can write  $h_3$  explicitly in terms of  $h_0, x_1, x_2, x_3$ :

$$h_1 = w_{\text{in}}x_1 + w_{\text{hh}}h_0,$$

$$h_2 = w_{\text{in}}x_2 + w_{\text{hh}}h_1 = w_{\text{in}}x_2 + w_{\text{hh}}(w_{\text{in}}x_1 + w_{\text{hh}}h_0),$$

$$h_3 = w_{\text{in}}x_3 + w_{\text{hh}}h_2 = w_{\text{in}}x_3 + w_{\text{hh}}[w_{\text{in}}x_2 + w_{\text{hh}}(w_{\text{in}}x_1 + w_{\text{hh}}h_0)].$$

Hence

$$h_3 = h_0w_{\text{hh}}^3 + w_{\text{in}}w_{\text{hh}}^2x_1 + w_{\text{in}}w_{\text{hh}}x_2 + w_{\text{in}}x_3.$$

Therefore

$$y = w_{\text{hy}}h_3 = w_{\text{hy}}(h_0w_{\text{hh}}^3 + w_{\text{in}}w_{\text{hh}}^2x_1 + w_{\text{in}}w_{\text{hh}}x_2 + w_{\text{in}}x_3).$$

Now:

$$\frac{\partial \ell}{\partial w_{\text{hy}}} = (y - y_r)h_3,$$

with  $h_3$  as above.

For  $w_{\text{in}}$  and  $w_{\text{hh}}$ , we first noted that

$$\frac{\partial \ell}{\partial h_3} = (y - y_r)w_{\text{hy}}, \quad \frac{\partial \ell}{\partial w_{\text{in}}} = \frac{\partial \ell}{\partial h_3} \frac{\partial h_3}{\partial w_{\text{in}}}, \quad \frac{\partial \ell}{\partial w_{\text{hh}}} = \frac{\partial \ell}{\partial h_3} \frac{\partial h_3}{\partial w_{\text{hh}}}.$$

And when we differentiate  $h_3$ :

$$\frac{\partial h_3}{\partial w_{\text{in}}} = w_{\text{hh}}^2 x_1 + w_{\text{hh}} x_2 + x_3,$$

$$\frac{\partial h_3}{\partial w_{\text{hh}}} = 3h_0 w_{\text{hh}}^2 + 2w_{\text{hh}} w_{\text{in}} x_1 + w_{\text{in}} x_2.$$

Therefore:

$$\boxed{\frac{\partial \ell}{\partial w_{\text{in}}} = (y - y_r) w_{\text{hy}} (w_{\text{hh}}^2 x_1 + w_{\text{hh}} x_2 + x_3)}$$

$$\boxed{\frac{\partial \ell}{\partial w_{\text{hh}}} = (y - y_r) w_{\text{hy}} (3h_0 w_{\text{hh}}^2 + 2w_{\text{hh}} w_{\text{in}} x_1 + w_{\text{in}} x_2)}$$

$$\boxed{\frac{\partial \ell}{\partial w_{\text{hy}}} = (y - y_r) (h_0 w_{\text{hh}}^3 + w_{\text{in}} w_{\text{hh}}^2 x_1 + w_{\text{in}} w_{\text{hh}} x_2 + w_{\text{in}} x_3)}$$

where  $y$  itself is

$$y = w_{\text{hy}} (h_0 w_{\text{hh}}^3 + w_{\text{in}} w_{\text{hh}}^2 x_1 + w_{\text{in}} w_{\text{hh}} x_2 + w_{\text{in}} x_3).$$

All expressions use only  $y_r, h_0, w_{\text{in}}, w_{\text{hh}}, w_{\text{hy}}, x_1, x_2, x_3$ .

**(d)** We already have from part b:

$$y = 0.01248.$$

Substitute  $x_1 = 0.3, x_2 = 0.4, x_3 = 0.2, h_0 = 0.5, w_{\text{in}} = w_{\text{hh}} = w_{\text{hy}} = 0.2, y_r = 0.4$  into the gradient formulas. This gives us

$$\frac{\partial \ell}{\partial w_{\text{in}}} \approx -0.022631168, \quad \frac{\partial \ell}{\partial w_{\text{hh}}} \approx -0.012710656, \quad \frac{\partial \ell}{\partial w_{\text{hy}}} \approx -0.024181248.$$

Using gradient descent

$$w^{\text{new}} = w^{\text{old}} - \eta \frac{\partial \ell}{\partial w}, \quad \eta = 0.15,$$

we obtain:

$$w_{\text{in}}^{\text{new}} = 0.2 - 0.15 \cdot (-0.022631168) \approx 0.2033947,$$

$$w_{\text{hh}}^{\text{new}} = 0.2 - 0.15 \cdot (-0.012710656) \approx 0.2019066,$$

$$w_{\text{hy}}^{\text{new}} = 0.2 - 0.15 \cdot (-0.024181248) \approx 0.2036272.$$

## Question 7

Each  $x_i$  is projected into:

$$q_i = W_Q x_i \quad (\text{query}), \quad k_i = W_K x_i \quad (\text{key}), \quad v_i = W_V x_i \quad (\text{value}),$$

where  $W_Q, W_K, W_V$  are the learned matrices.

To compute the output representation for  $x_1$ :

1. Compute the query for position 1:

$$q_1 = W_Q x_1.$$

2. Compute keys and values for all positions  $j = 1, 2, 3$ :

$$k_j = W_K x_j, \quad v_j = W_V x_j.$$

3. Compute attention scores between  $q_1$  and each  $k_j$ :

$$e_{1j} = \frac{q_1^\top k_j}{\sqrt{d_k}},$$

where  $d_k$  is the dimensionality of the keys.

4. Normalize scores with softmax to obtain attention weights:

$$\alpha_{1j} = \frac{\exp(e_{1j})}{\sum_{m=1}^3 \exp(e_{1m})}, \quad j = 1, 2, 3.$$

5. The output representation for position 1 is the weighted sum of values:

$$z_1 = \sum_{j=1}^3 \alpha_{1j} v_j.$$

Here:

- Queries  $q_1$  ask “what am I looking for?” from position 1.
- Keys  $k_j$  encode “what content do I have?” at each position  $j$ .
- Values  $v_j$  carry the information that will be mixed together.

The self-attention mechanism lets  $z_1$  integrate information from  $x_1, x_2, x_3$  according to the learned attention weights  $\alpha_{1j}$ .